



# KNN-Based Image Search with Elasticsearch

Leveraging vector embeddings for efficient image similarity search

NAVIN BALAJI RK

# Introduction

- Problem: Searching for similar images is complex using traditional methods.
- Goal: Use image embeddings + KNN to search similar images efficiently.
- Tools: Javascript, CLIP, Elasticsearch, KNN vector search

The background is a dark navy blue. It features several abstract geometric elements: a large teal semi-circle in the top right, a smaller orange semi-circle below it, a teal rectangle to the right of the top semi-circle, an orange rectangle below that, a teal semi-circle in the bottom right, and a teal rectangle to its left. There are also some faint, lighter blue circular patterns in the background.

# What is Vector ?

Vector embeddings are a way to convert words and sentences and other data like images into numbers that capture their meaning and relationship




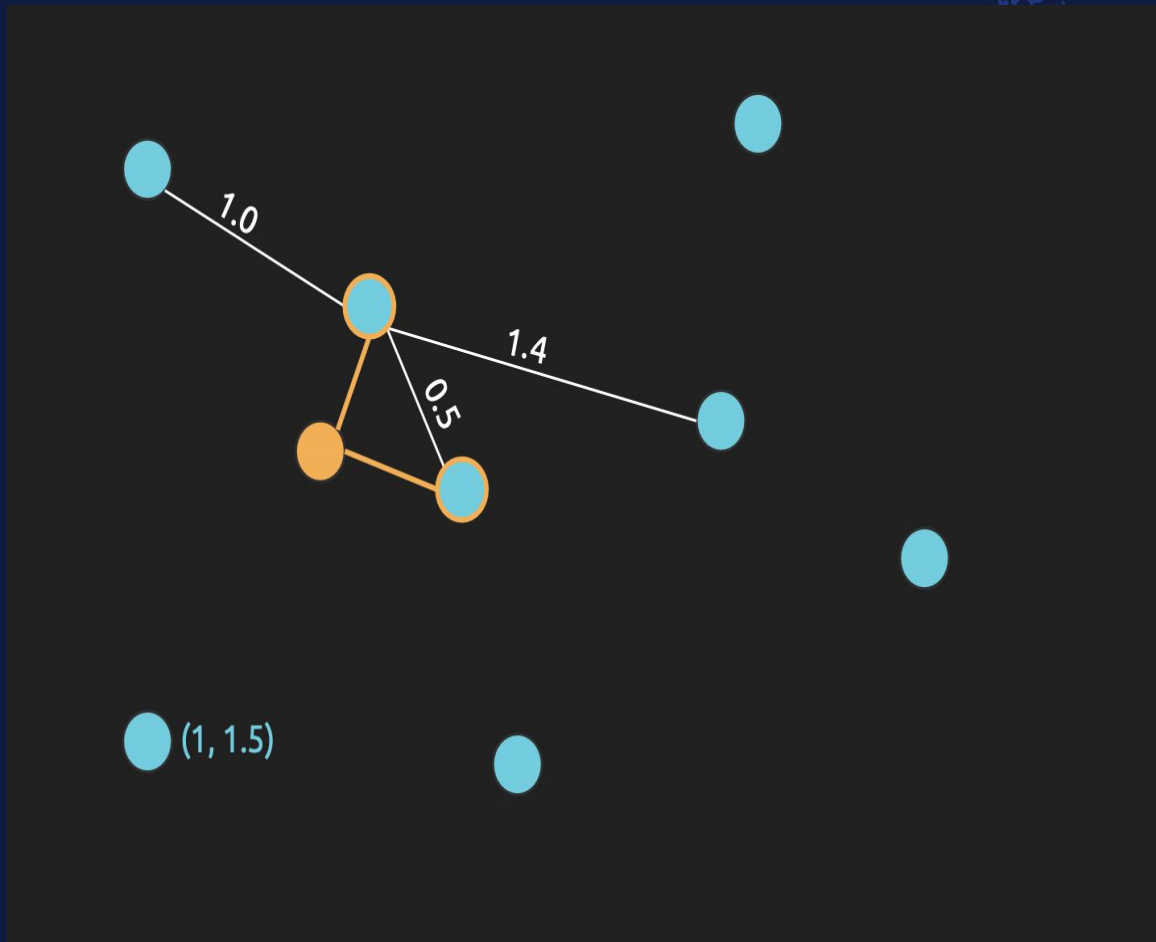
# Different type of Vectors



1. Dense Vector
2. Sparse Vector

What is KNN?





Fluffiness and Lightness of Animals Dataset





# Image Vectorization

- Use pre-trained models like OpenAI CLIP or Xenova/clip-vit-base-patch32
- Convert images to 512-D

JS elastic.js

```
import { Client } from "@elastic/elasticsearch";

const client = new Client({
  node: process.env.ELASTIC_NODE,
  auth: {
    username: process.env.ELASTIC_USERNAME,
    password: process.env.ELASTIC_PASSWORD
  }
});
```



CodelImage



JS elastic.js

```
await client.indices.create({
  index: INDEX_NAME,
  body: {
    mappings: {
      properties: {
        image_vector: {
          type: "dense_vector",
          dims: 512,
          index: true,
          similarity: "cosine"
        },
        filename: { type: "keyword" },
        uploaded_at: { type: "date" }
      }
    }
  }
});
```



CodelImage

elastic.js

```
import { pipeline } from "@xenova/transformers";

const extractor = await pipeline(
  "image-feature-extraction",
  "Xenova/clip-vit-base-patch32"
);

const embeddings = await extractor(imagePath, {
  pooling: "mean",
  normalize: true
});

return Array.from(embeddings.data);
```



CodelImage



elastic.js

```
await client.index({  
  index: INDEX_NAME,  
  document: {  
    image_vector: vector,  
    filename,  
    uploaded_at: new Date()  
  }  
});
```



CodelImage



elastic.js

```
const response = await client.search({  
  index: INDEX_NAME,  
  knn: {  
    field: "image_vector",  
    query_vector: vector,  
    k: 5,  
    num_candidates: 50  
  },  
  _source: ["filename", "uploaded_at"]  
});
```



CodelImage

# Elasticsearch Plugin for Nearest Neighbor Search

The `knn` search option accepts a number of parameters that configure the search:

- `field`: the field in the index to search. The field must have a `dense_vector` type.
- `query_vector`: the embedding to search for. This should be an embedding generated from the search text.
- `num_candidates`: the number of candidate documents to consider from each shard. Elasticsearch retrieves this many candidates from each shard, combines them into a single list and then finds the closest "k" to return as results.
- `k`: the number of results to return. This number has a direct effect on performance, so it should be kept as small as possible. The value passed in this option must be less than `num_candidates`.

•

# Resources

- [Dense vector field type | Elasticsearch Guide \[8.17\] | Elastic](#)
- [k-nearest neighbor \(kNN\) search | Elasticsearch Guide \[8.17\] | Elastic](#)
- [Knn query | Elasticsearch Guide \[8.17\] | Elastic](#)
- <https://www.elastic.co/search-labs/blog/implementing-image-search-with-elasticsearch>
- <https://huggingface.co/Xenova/clip-vit-base-patch32>
- <https://alexklibisz.github.io/elastiknn>





Questions?



# Thank You

More about me

